

CS 630 - 002

Operating Systems Design

Lecture 2

OS Designs

Jing Li

jingli@njit.edu

GITC 4106



New Jersey Institute of Technology

In This Lecture

- Different Operating Systems Designs
- Administrivia
 - ❑ Assignment 1 will be released on Feb 5 and due on Feb 19 23:55
- OS Examples:
 - ❑ Unix, Linux, Window, Mac
- Background
 - ❑ How to gradually add abstraction to multiplex resources

Slides courtesy of Hung Daochuan, Chris Gill, David Ferry, Tarek Abdelzaher, Ion Stoica, John Kubiawicz, Peter Dennings, Anthony Joseph, Jonathan Ragan-Kelley, Peter Troger.

Recap



New Jersey Institute of Technology

Abstraction, Arbitration & Protection

- OS offers abstractions, arbitration and protection

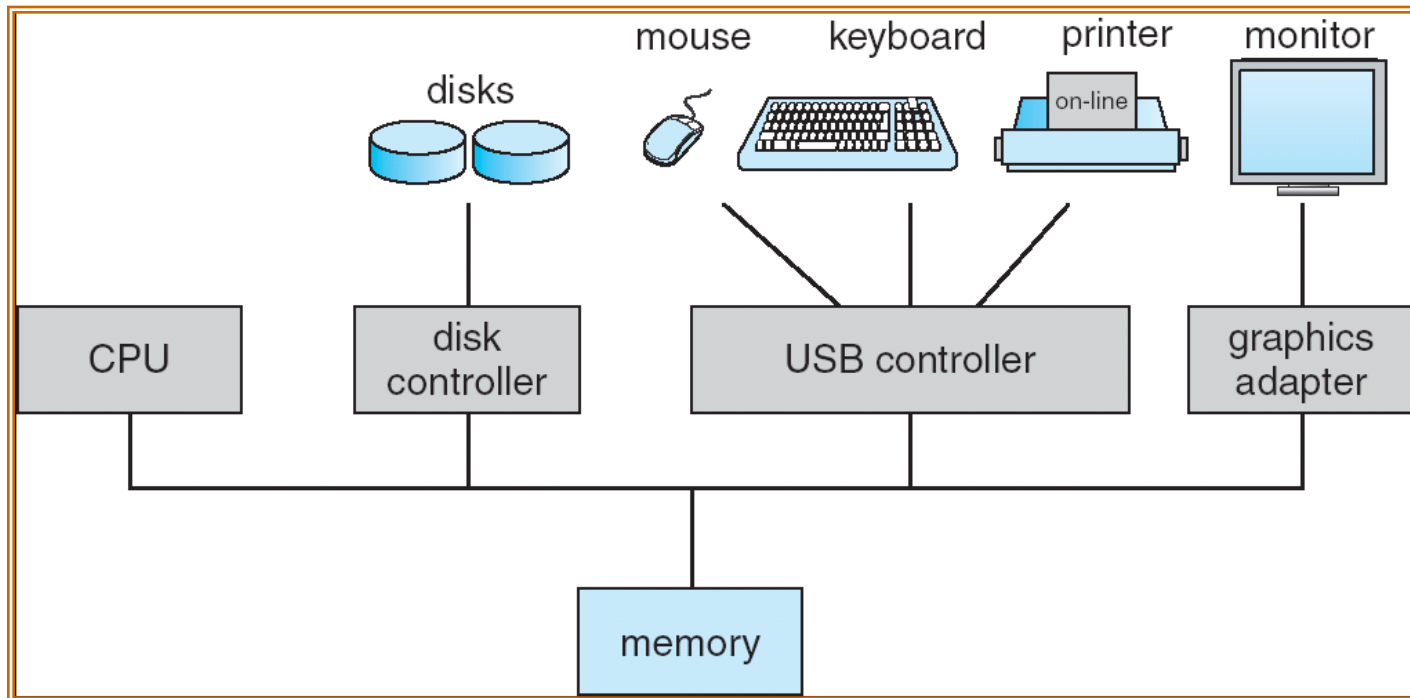
Application

Operating System

Hardware

Virtual Machine Interface

Physical Machine Interface




Common Operating System Components

- Process Management:
 - ❑ "process" is a program in execution. OS responsible for creation, deletion, scheduling, communication
- Main-memory Management: allocation, protection
- File Management:
 - ❑ creation, deletion, directory structures, mapping files to hardware
- I/O System Management: device driver interface, buffering
- Secondary Storage Management
 - ❑ free space management, storage allocation, disk scheduling
- Networking: another device to manage - high-speed information flow
- Protection System: specification and enforcement of access controls
- Command-interpreter:
 - ❑ e.g. UNIX command line or MS-DOS prompt. A windowing system is just a way to issue the same commands without knowing what they are.


Evolution of CPU utilizations



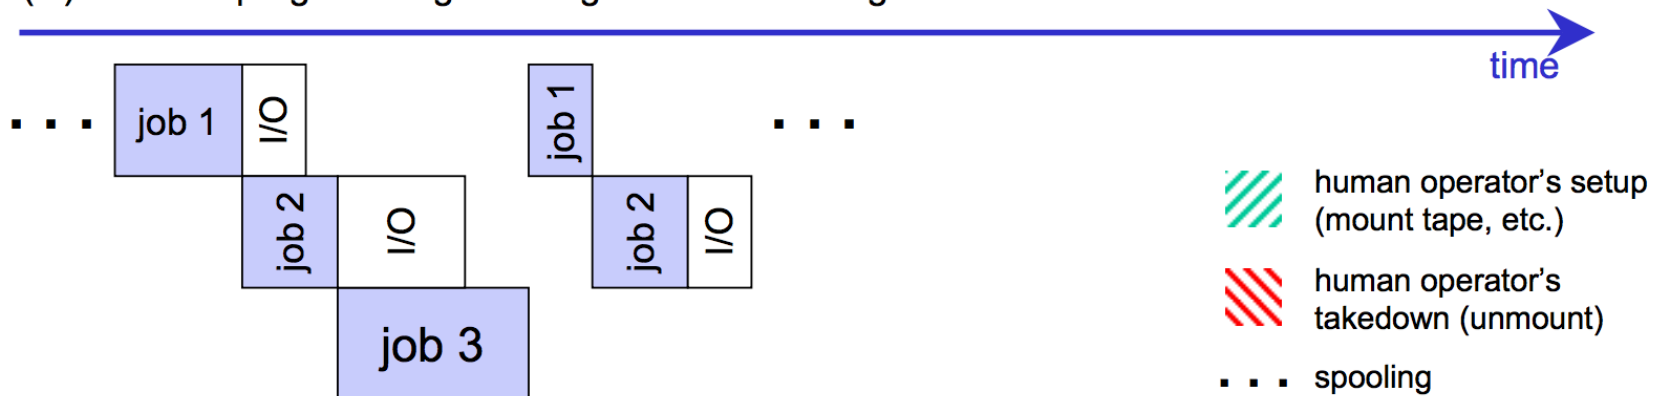
(a) serial uniprogramming



(b) batch uniprogramming



(b') batch uniprogramming showing actual CPU usage and I/O wait

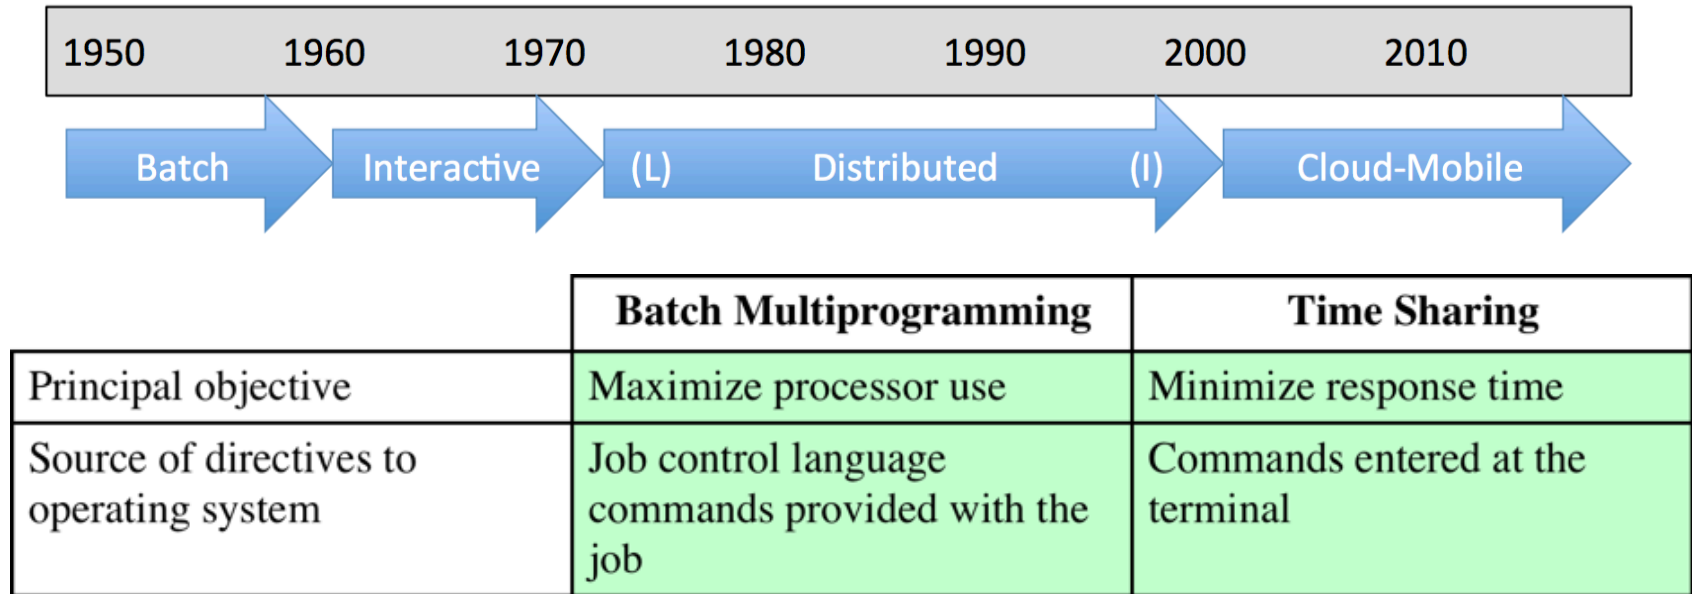


(c) multiprogramming

Evolution of CPU utilization

Eras of Operating Systems

➤ Time sharing systems (multitasking)



- This is done by switching user tasks or **processes** transparently. We want each interactive user to get a turn on the CPU quickly - good **response time**. Need to switch among processes quickly - **context switching**.

Different Operating Systems Designs



New Jersey Institute of Technology

Simple Structure

- All aspects of the OS linked together in one binary
 - ❑ APIs not carefully designed (and/or lots of global variables)
 - ❑ Interfaces and levels of functionality not well separated
 - ❑ No address protection

- Example: MS-DOS

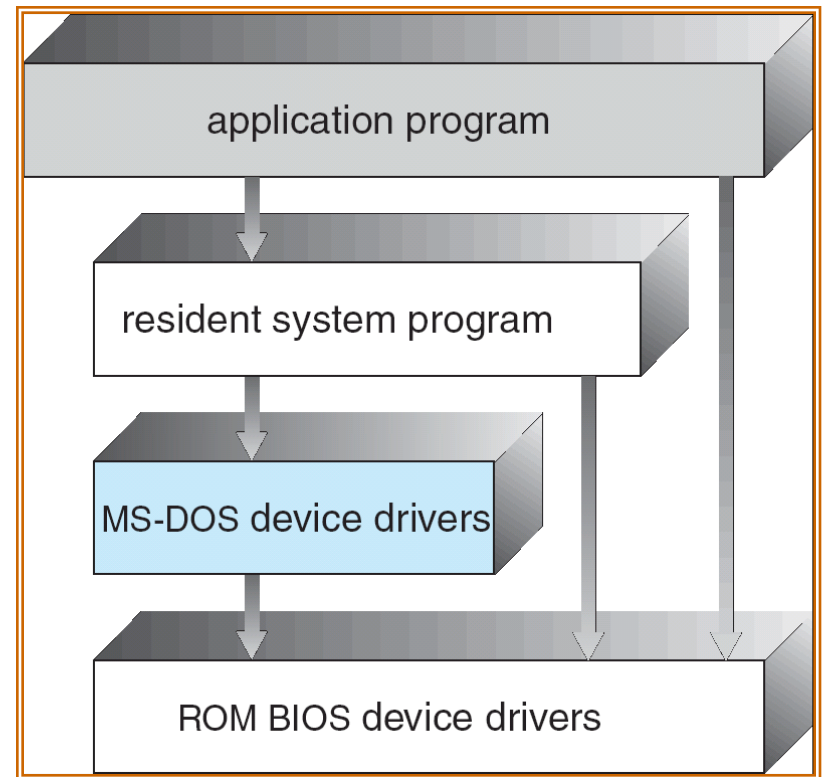
- ❑ provide the most functionality in the least space
 - ❑ Made sense in early days of personal computers with limited processors (e.g. 6502)

- Advantages?

- ❑ Low memory footprint

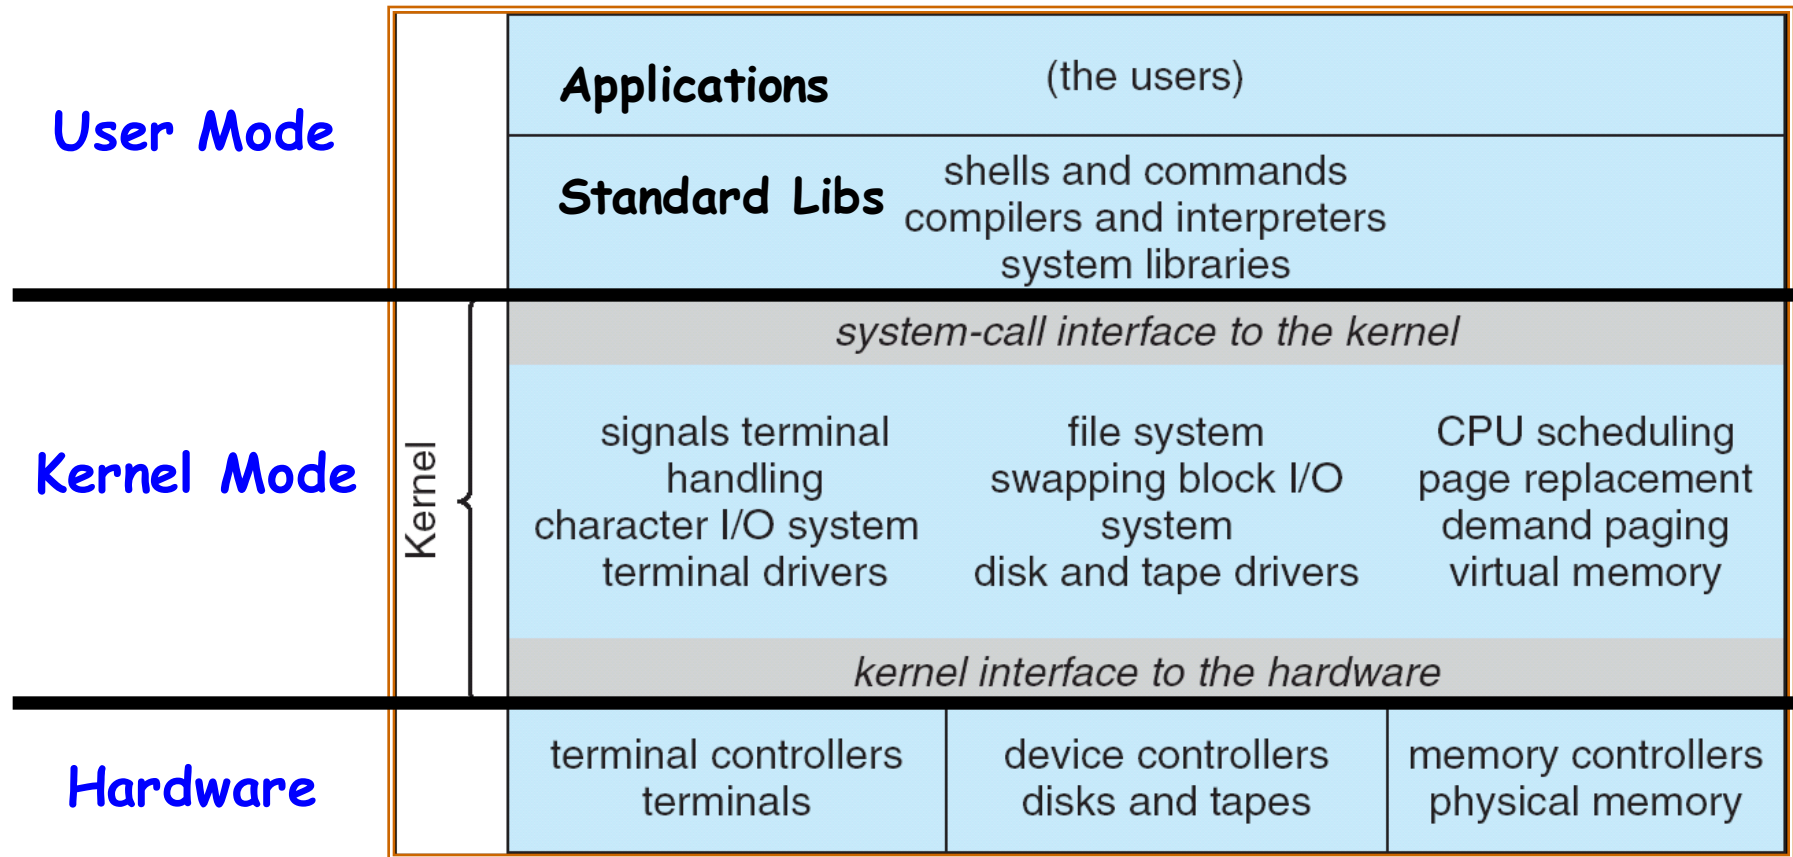
- Disadvantages?

- ❑ Very fragile, no enforcement of structure/boundaries




Monolithic Structure: UNIX System Structure

➤ Two-Layered Structure: User vs Kernel



Layered Operating System

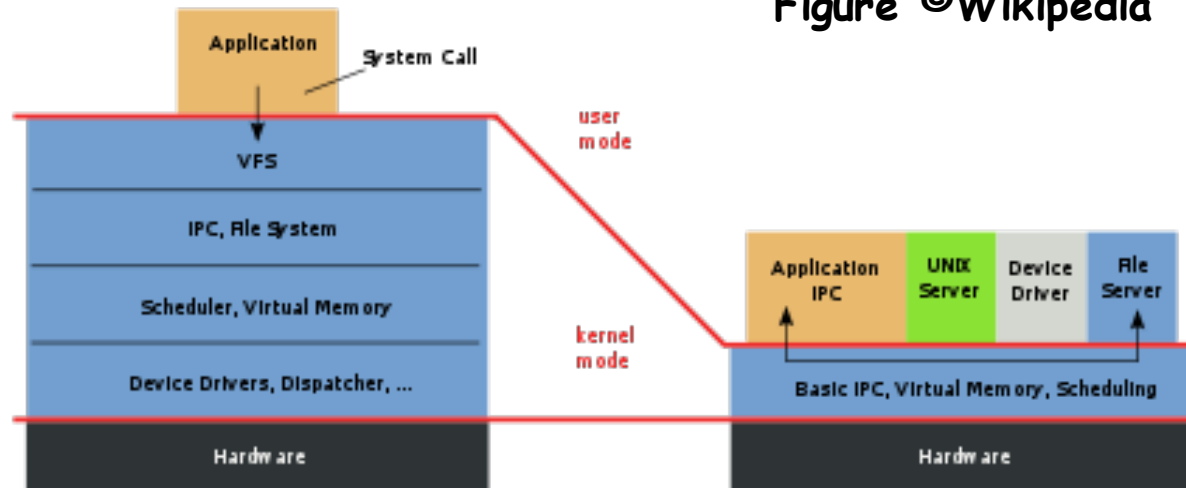
- Operating system is divided into many layers (levels)
 - ❑ Each built on top of lower layers
 - ❑ Bottom layer (layer 0) is hardware
 - ❑ Highest layer (layer N) is the user interface
- Each layer uses functions (operations) and services of only lower-level layers
- Machine-dependent vs. independent layers



Microkernel Structure

Figure ©Wikipedia

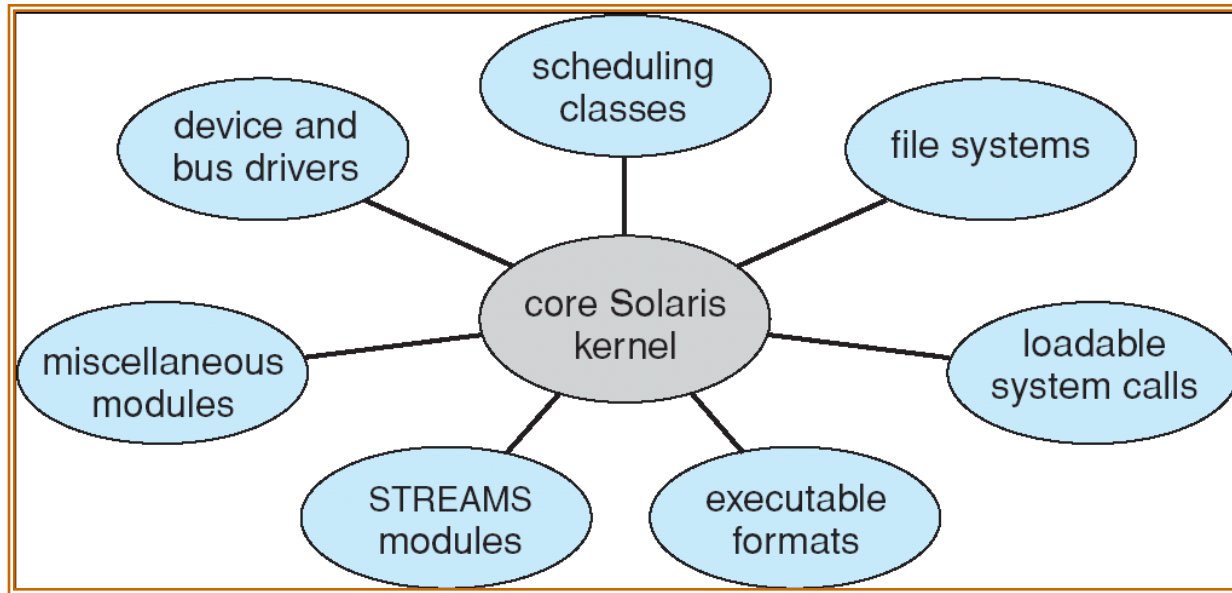
Monolithic
Kernel



Microkernel

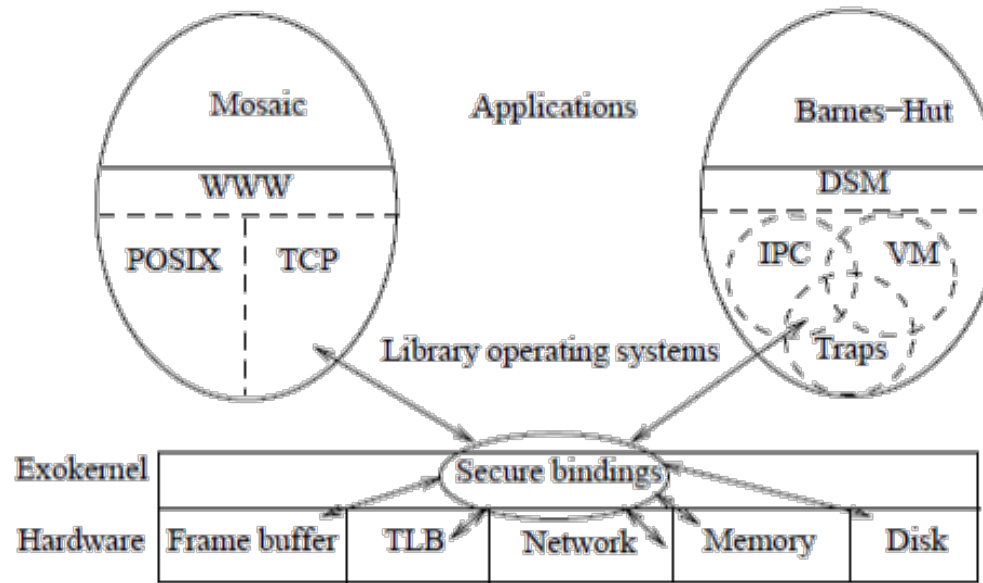
- Moves functionality from the kernel into “user” space
 - ❑ Small core OS running at kernel level
 - ❑ OS Services built from many independent user-level processes
 - ❑ Communication between modules with message passing
- Benefits:
 - ❑ Easier to extend a microkernel
 - ❑ Easier to port OS to new architectures
 - ❑ More reliable (less code is running in kernel mode)
 - ❑ Fault Isolation (parts of kernel protected from other parts)
 - ❑ More secure
- Detriments:
 - ❑ Performance overhead can be severe for naïve implementation

Modules-based Structure



- Most modern operating systems implement modules
 - ❑ Uses object-oriented approach
- Each core component is separate
 - ❑ Each talks to the others over known interfaces
 - ❑ Each is loadable as needed within the kernel
- Overall, similar to layers but with more flexible
 - ❑ May or may not utilize hardware enforcement

ExoKernel: Separate Protection from Management



- Thin layer exports hardware resources directly to users
 - ❑ As little abstraction as possible
 - ❑ Secure Protection and Multiplexing of resources
- LibraryOS: traditional OS functionality at User-Level
 - ❑ Customize resource management for every application
 - ❑ Is this a practical approach?
- Very low-level abstraction layer
 - ❑ Need extremely specialized skills to develop LibraryOS

Different Operating Systems Designs

- Simple:
 - ❑ Only one or two levels of code
- Layered:
 - ❑ Lower levels independent of upper levels
- Microkernel:
 - ❑ OS built from many user-level processes
- Modular:
 - ❑ Core kernel with Dynamically loadable modules
- ExoKernel:
 - ❑ Separate protection from management of resources
- Cell-based OS:
 - ❑ Two-level scheduling of resources

Videos

[History of Apple \(OS & Others\)](#)

[History of Mac OS \(UI & Features\)](#)



New Jersey Institute of Technology

Administrivia

**Office hour update: Shaoze Fan,
Wednesday from 10:00 AM to 11:30
AM GITC 4325.**



New Jersey Institute of Technology

Labs

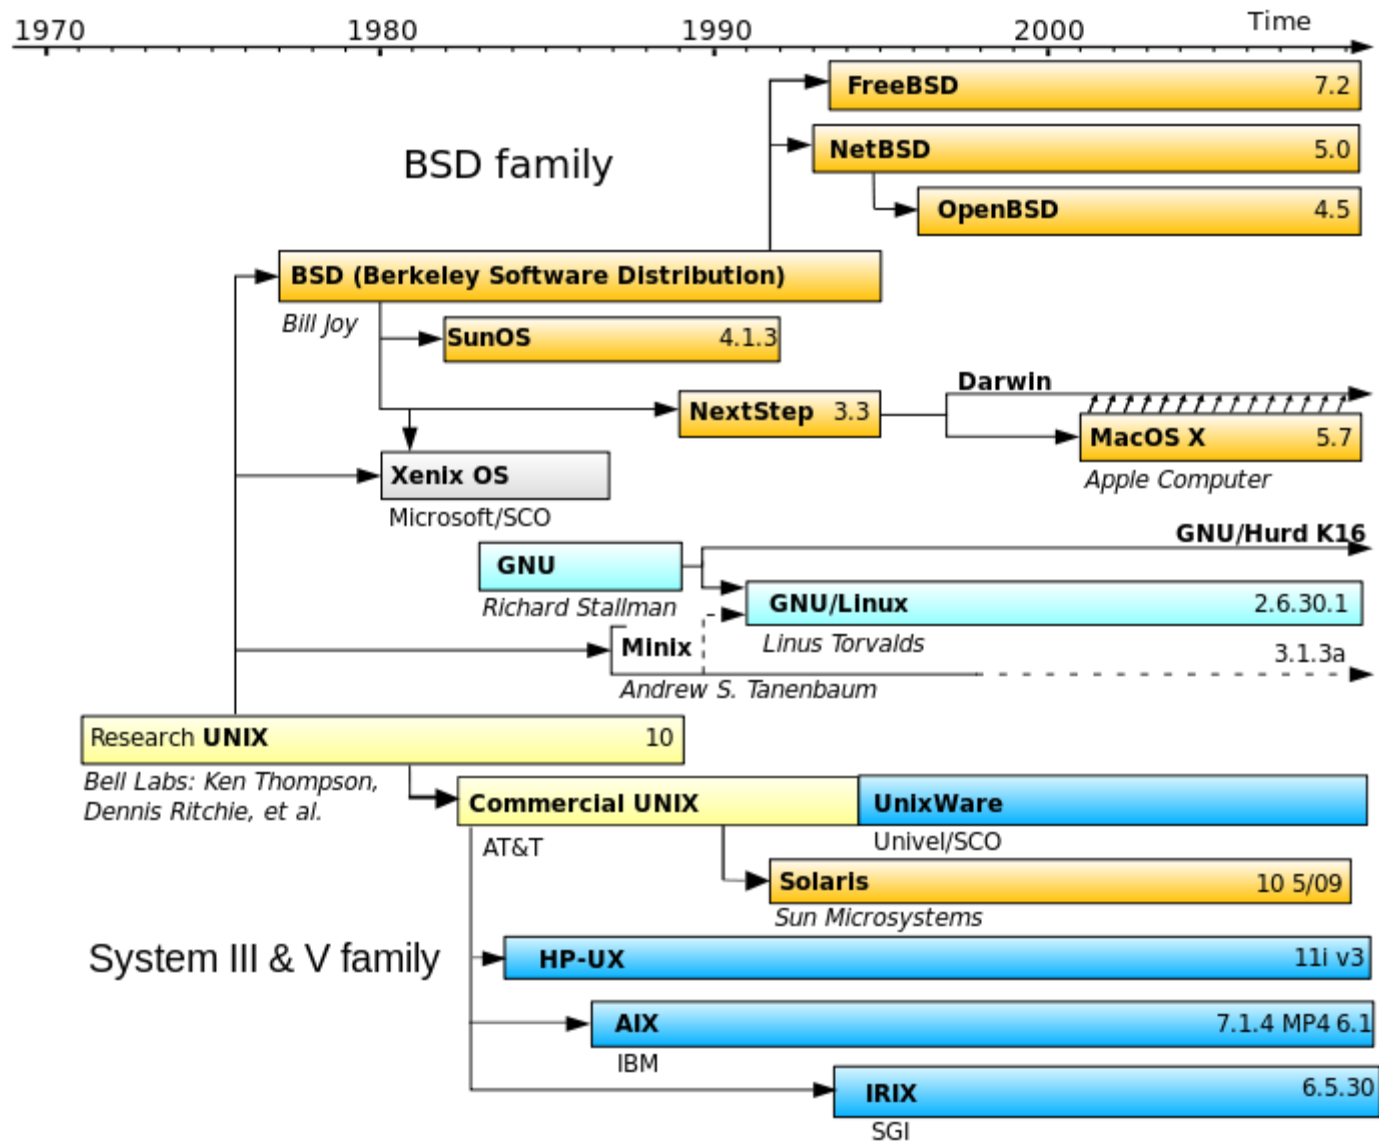
- Remember concepts is different from using/implementing
- If time allows, we will have in-class demo/lab time
- Form your team ASAP:
 - ❑ Each lab will be completed in a team of up to 3 students, and teams may be different for each lab.
- Lab I will be released on Feb 5 and **due on Feb 19 23:55**
 - ❑ Your first trial of forming your group
 - ❑ Ease you into programming in Linux
 - ❑ Bases for a later Lab – writing a shell program

Example Operating Systems



New Jersey Institute of Technology

History of Unix



Videos

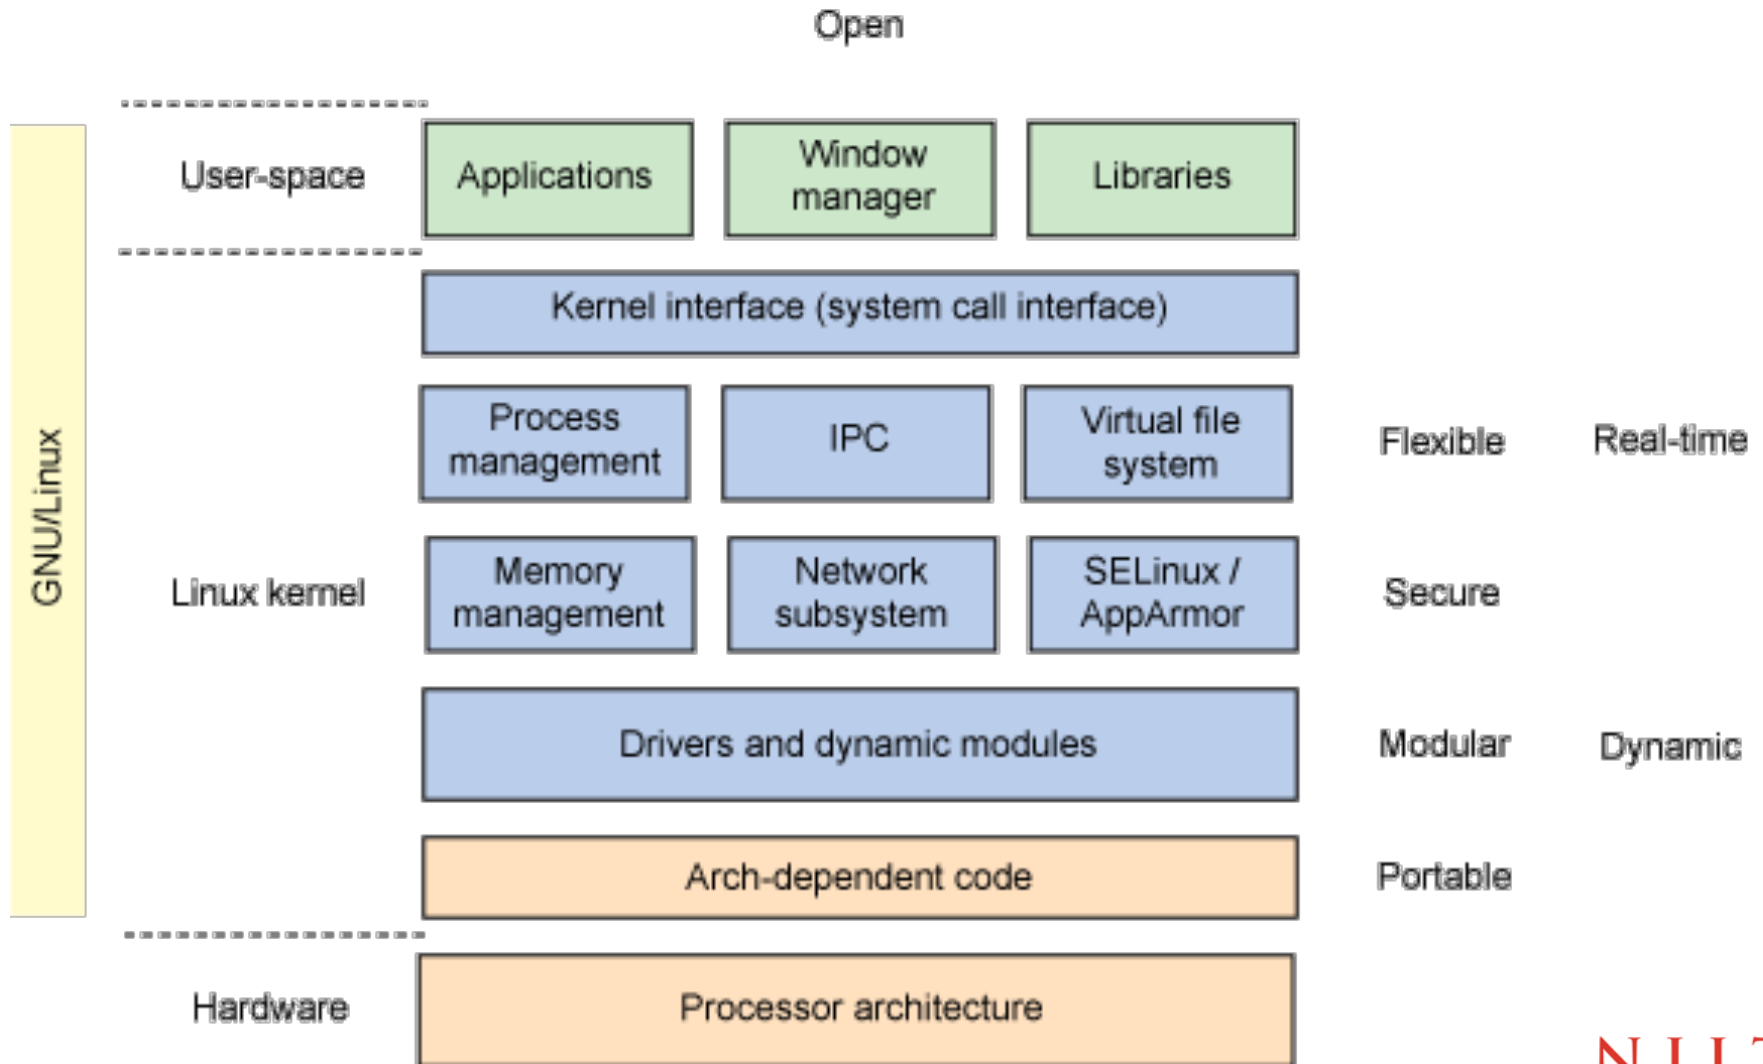
[Unix to Linux \(Part 1: first 3min & last 1min\)](#)
[\(Part 2: first 3min\)](#)

[Linux: A Short Documentary](#)




New Jersey Institute of Technology

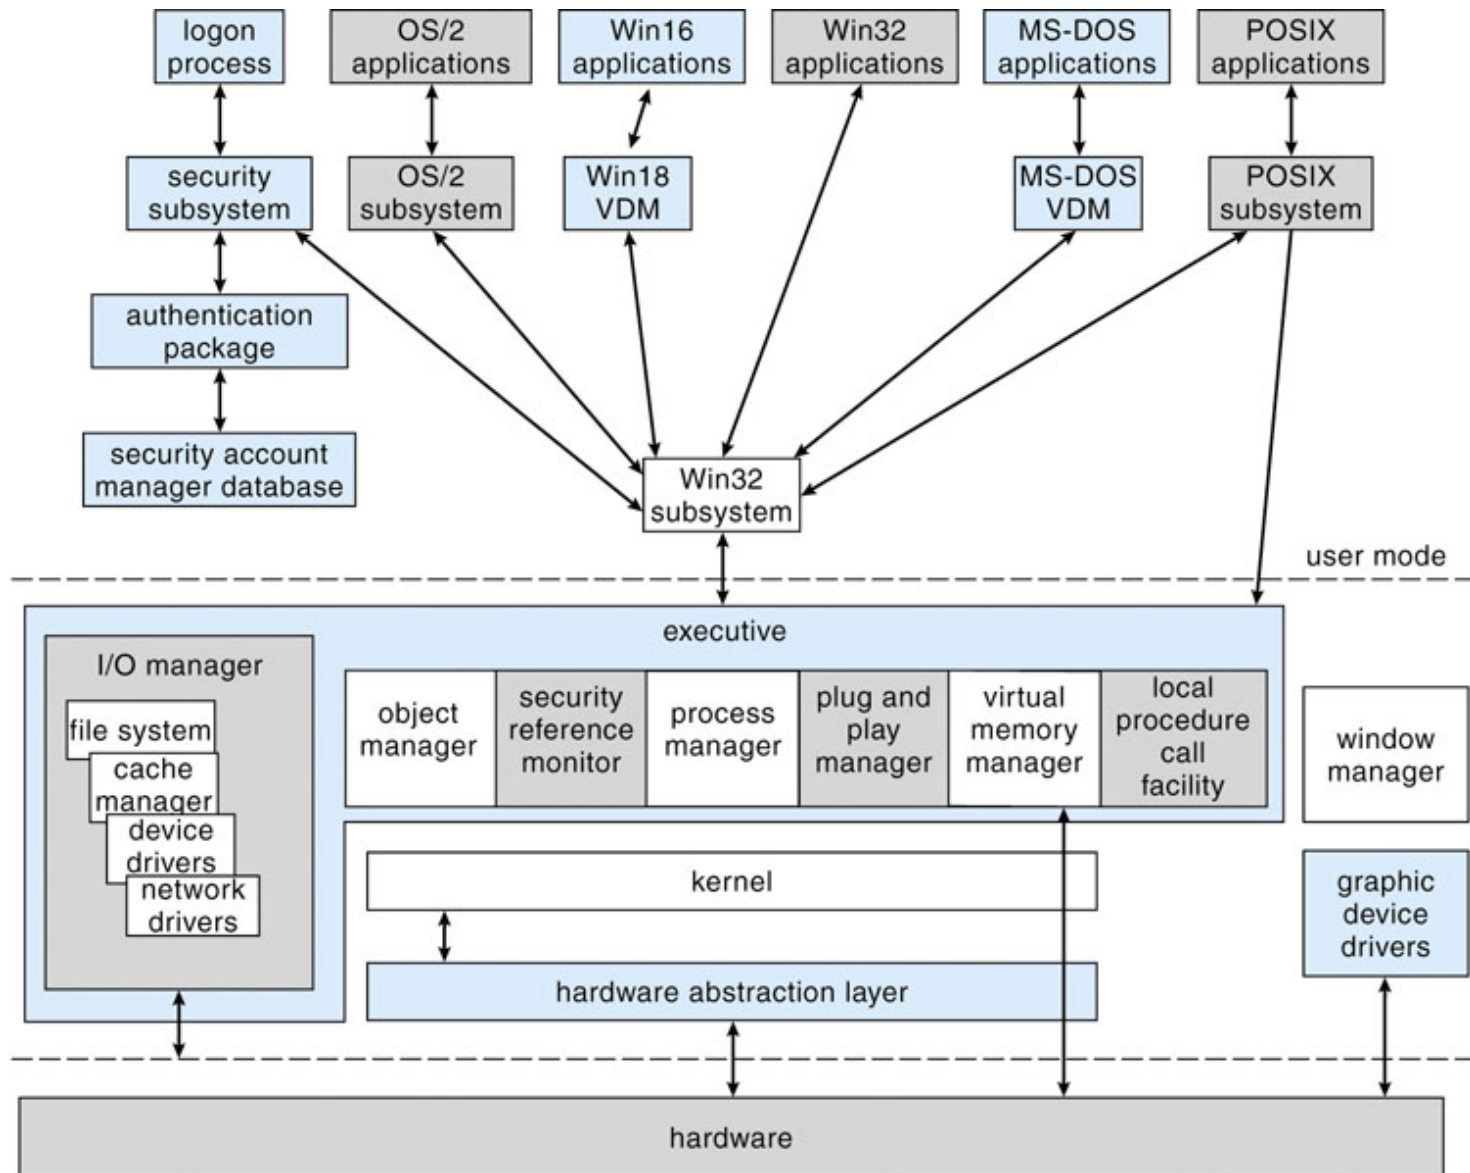
Linux Structure (Simplified Figure)




➤ http://www.makelinux.net/kernel_map/



Microsoft Windows Structure



History of Windows



Videos

The Evolution of Windows® – From
Windows® 1.0 to Windows® 10



New Jersey Institute of Technology

Background



New Jersey Institute of Technology

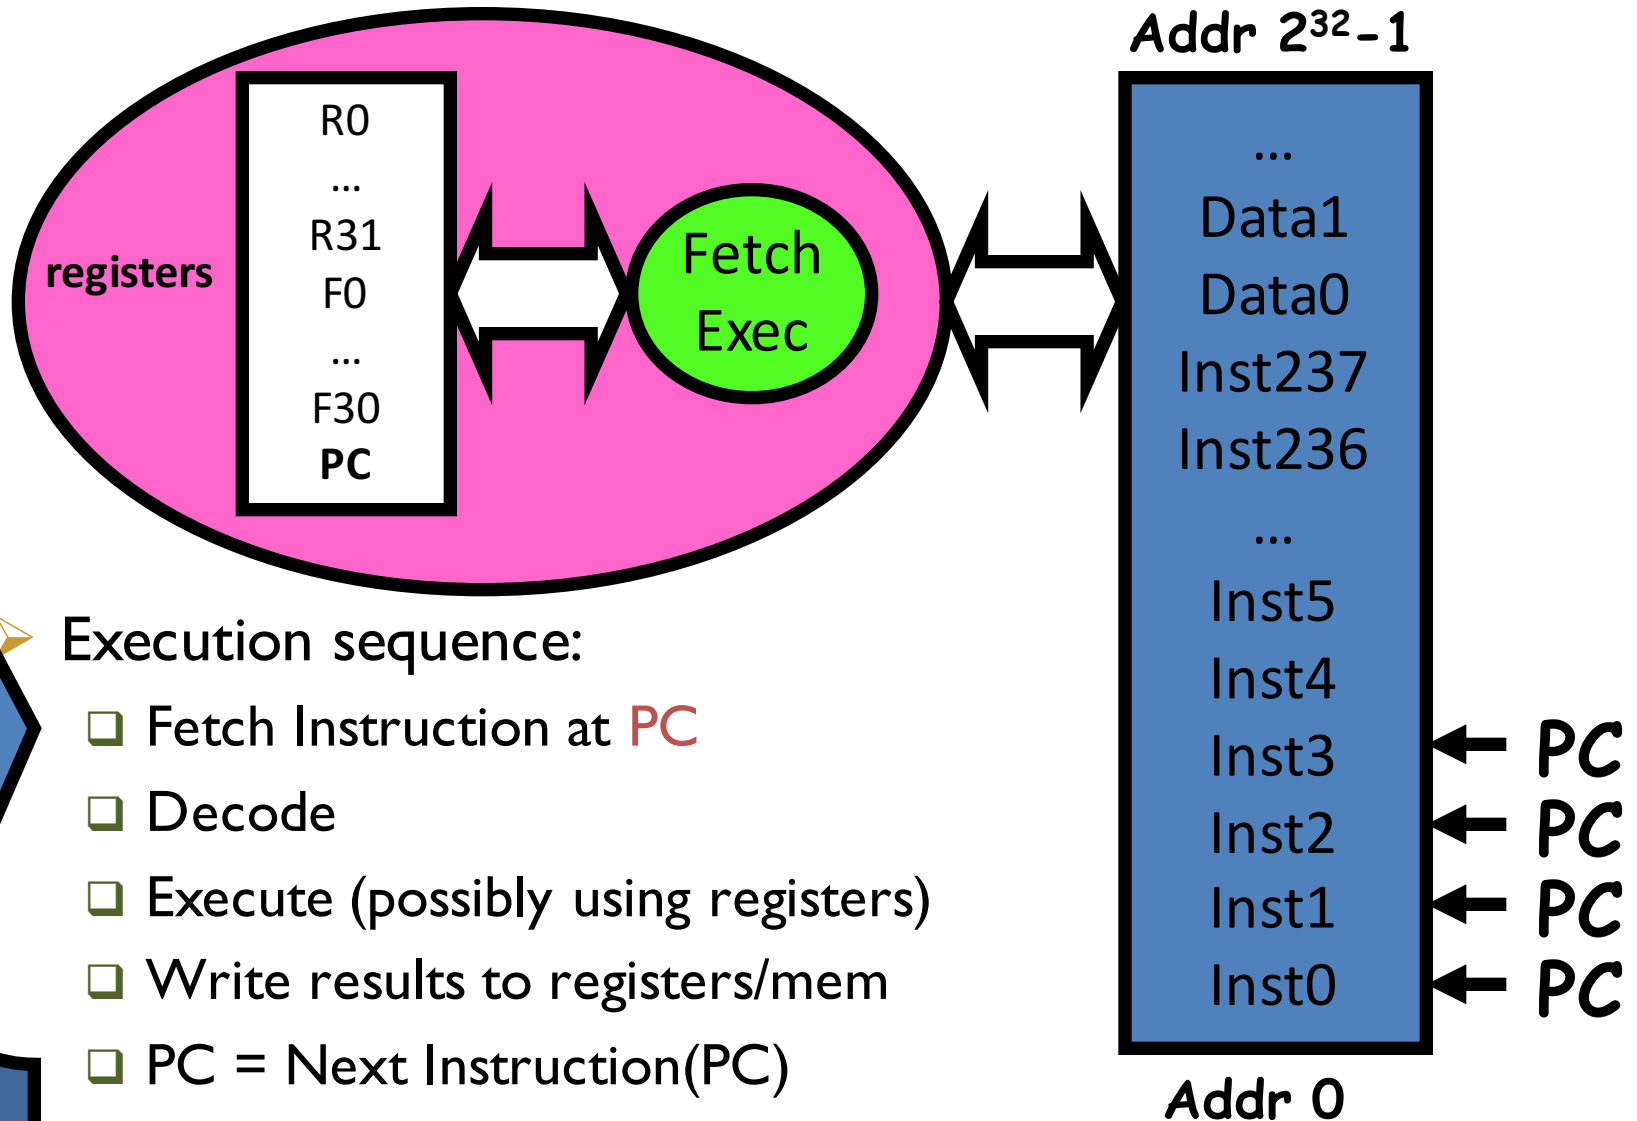
Concurrency

- “Thread” of execution
 - ❑ Independent Fetch/Decode/Execute loop
 - ❑ Operating in some Address space
- Uniprogramming: *one thread at a time*
 - ❑ MS/DOS, early Macintosh, Batch processing
 - ❑ Easier for operating system builder
 - ❑ Get rid of concurrency by defining it away
 - ❑ Does this make sense for personal computers?
- Multiprogramming: *more than one thread at a time*
 - ❑ Multics, UNIX/Linux, OS/2, Windows NT/2000/XP, Mac OS X
 - ❑ Often called “multitasking”, but multitasking has other meanings
- Multicore \Rightarrow Multiprogramming, right?


The Basic Problem of Concurrency

- The basic problem of concurrency involves resources:
 - ❑ Hardware (simple): single CPU, single DRAM, single I/O devices
 - ❑ Multiprogramming API: users think they have exclusive access to shared resources
- OS Has to coordinate all activity
 - ❑ Multiple users, I/O interrupts, ...
 - ❑ How can it keep all these things straight?
- Basic Idea: Use abstraction
 - ❑ Decompose hard problem into simpler ones
 - ❑ Abstract the notion of an executing program
 - ❑ Then, worry about multiplexing these abstract machines

Recall: What happens during program execution?



How to give the illusion of multiple processors?



- Assume a single processor. How do we provide the illusion of multiple processors?
 - ❑ Multiplex in time!
- Each virtual “CPU” needs a structure to hold:
 - ❑ Program Counter (PC), Stack Pointer (SP)
 - ❑ Registers (Integer, Floating point, others...?)
 - ❑ **Call result a “Thread” for now...**
- How to switch from one CPU to the next?
 - ❑ Save PC, SP, and registers in current state block
 - ❑ Load PC, SP, and registers from new state block
- What triggers switch?
 - ❑ Timer, voluntary yield, I/O, other things

What needs to be saved in Modern X86?

64-bit Register Set

General Purpose Registers (GPRs)

	RAX
	RBX
	RCX
	RDX
	RBP
	RSI
	RDI
	RSP
	R8
	R9
	R10
	R11
	R12
	R13
	R14
	R15

63 0

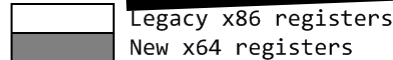

80-bit floating point
and 64-bit MMX registers
(overlaid)

	MMX Part
	FPR0/MMX0
	FPR1/MMX1
	FPR2/MMX2
	FPR3/MMX3
	FPR4/MMX4
	FPR5/MMX5
	FPR6/MMX6
	FPR7/MMX7



79 63 0

Also: 6 segment registers, control, status, debug, more

Address Space



Instruction Pointer/Flags



128-bit XMM Registers

	XMM0
	XMM1
	XMM2
	XMM3
	XMM4
	XMM5
	XMM6
	XMM7
	XMM8
	XMM9
	XMM10
	XMM11
	XMM12
	XMM13
	XMM14
	XMM15

127

0

Traditional 32-bit subset

General-Purpose Registers

31	16	15	8	7	0	16-bit	32-bit
			AH		AL	AX	EAX
			BH		BL	BX	EBX
			CH		CL	CX	ECX
			DH		DL	DX	EDX
			BP				EBP
			SI				ESI
			DI				EDI
			SP				ESP

EFLAGS Register

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0																						
0	0	0	0	0	0	0	0	0	0	0	D	I	V	I	P	A	C	V	M	R	F	0	0	N	T	0	I	O	P	L	0	F	F	0	F	F	I	F	T	F	S	F	S	F	0	A	F	0	P	F	1	C	F

- X ID Flag (ID)
- X Virtual Interrupt Pending (VIP)
- X Virtual Interrupt Flag (VIF)
- X Alignment Check (AC)
- X Virtual-8086 Mode (VM)
- X Resume Flag (RF)
- X Nested Task (NT)
- X I/O Privilege Level (IOPL)
- S Overflow Flag (OF)
- C Direction Flag (DF)
- X Interrupt Enable Flag (IF)
- X Trap Flag (TF)
- S Sign Flag (SF)
- S Zero Flag (ZF)
- S Auxiliary Carry Flag (AF)
- S Parity Flag (PF)
- S Carry Flag (CF)

S Indicates a Status Flag
C Indicates a Control Flag
X Indicates a System Flag

Property of the simple multiprogramming technique

- All virtual CPUs share same non-CPU resources
 - ❑ I/O devices the same
 - ❑ Memory the same
- Consequence of sharing:
 - ❑ Each thread can access the data of every other thread
(good for sharing, bad for protection)
 - ❑ Threads can share instructions
(good for sharing, bad for protection)
 - ❑ Can threads overwrite OS functions?
- This (unprotected) model common in:
 - ❑ Embedded applications
 - ❑ Windows 3.1/Machintosh (switch only with yield)
 - ❑ Windows 95 (switch with both yield and timer)

How to protect threads from one another?

Need three important things:

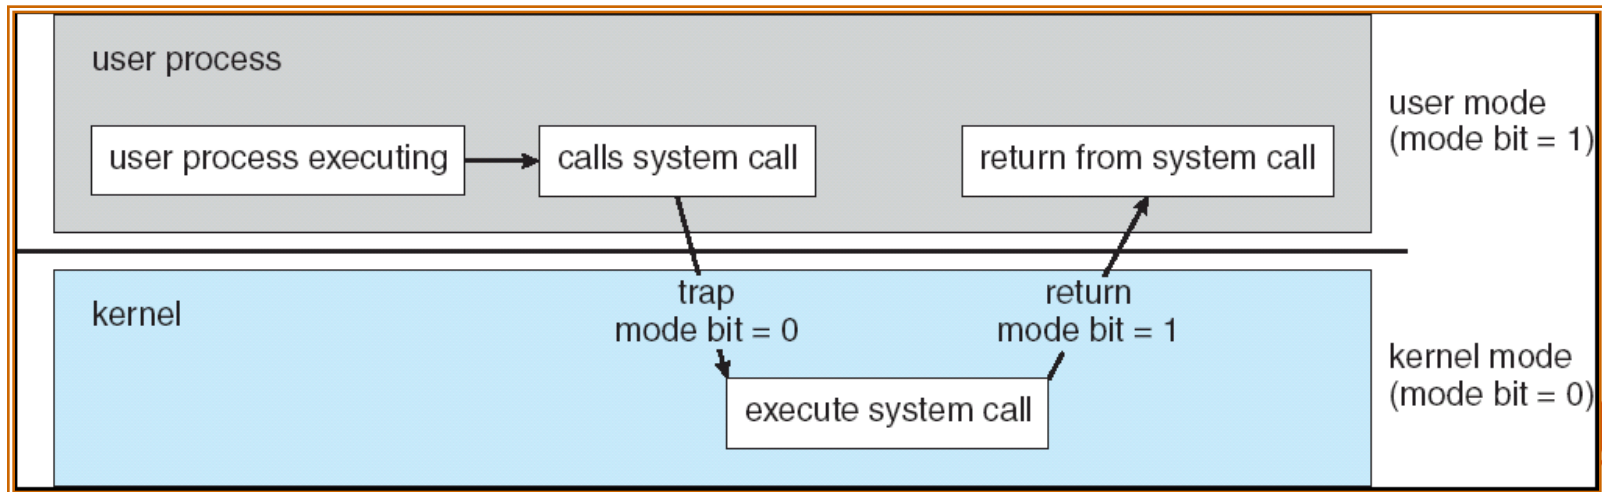
- Protection of memory
 - ❑ Every task does not have access to all memory
- Protection of I/O devices
 - ❑ Every task does not have access to every device
- Protection of Access to Processor:
Preemptive switching from task to task
 - ❑ Use of timer
 - ❑ Must not be possible to disable timer from user code

Protecting Threads from Each Other

- Problem: Run multiple applications in such a way that they are protected from one another
- Goal:
 - ❑ Keep User Programs from Crashing OS
 - ❑ Keep User Programs from Crashing each other
 - ❑ [Keep Parts of OS from crashing other parts?]
- (Some of the required) Mechanisms:
 - ❑ Address Translation
 - ❑ Dual Mode Operation
- Simple Policy:
 - ❑ Programs are not allowed to read/write memory of other Programs or of Operating System

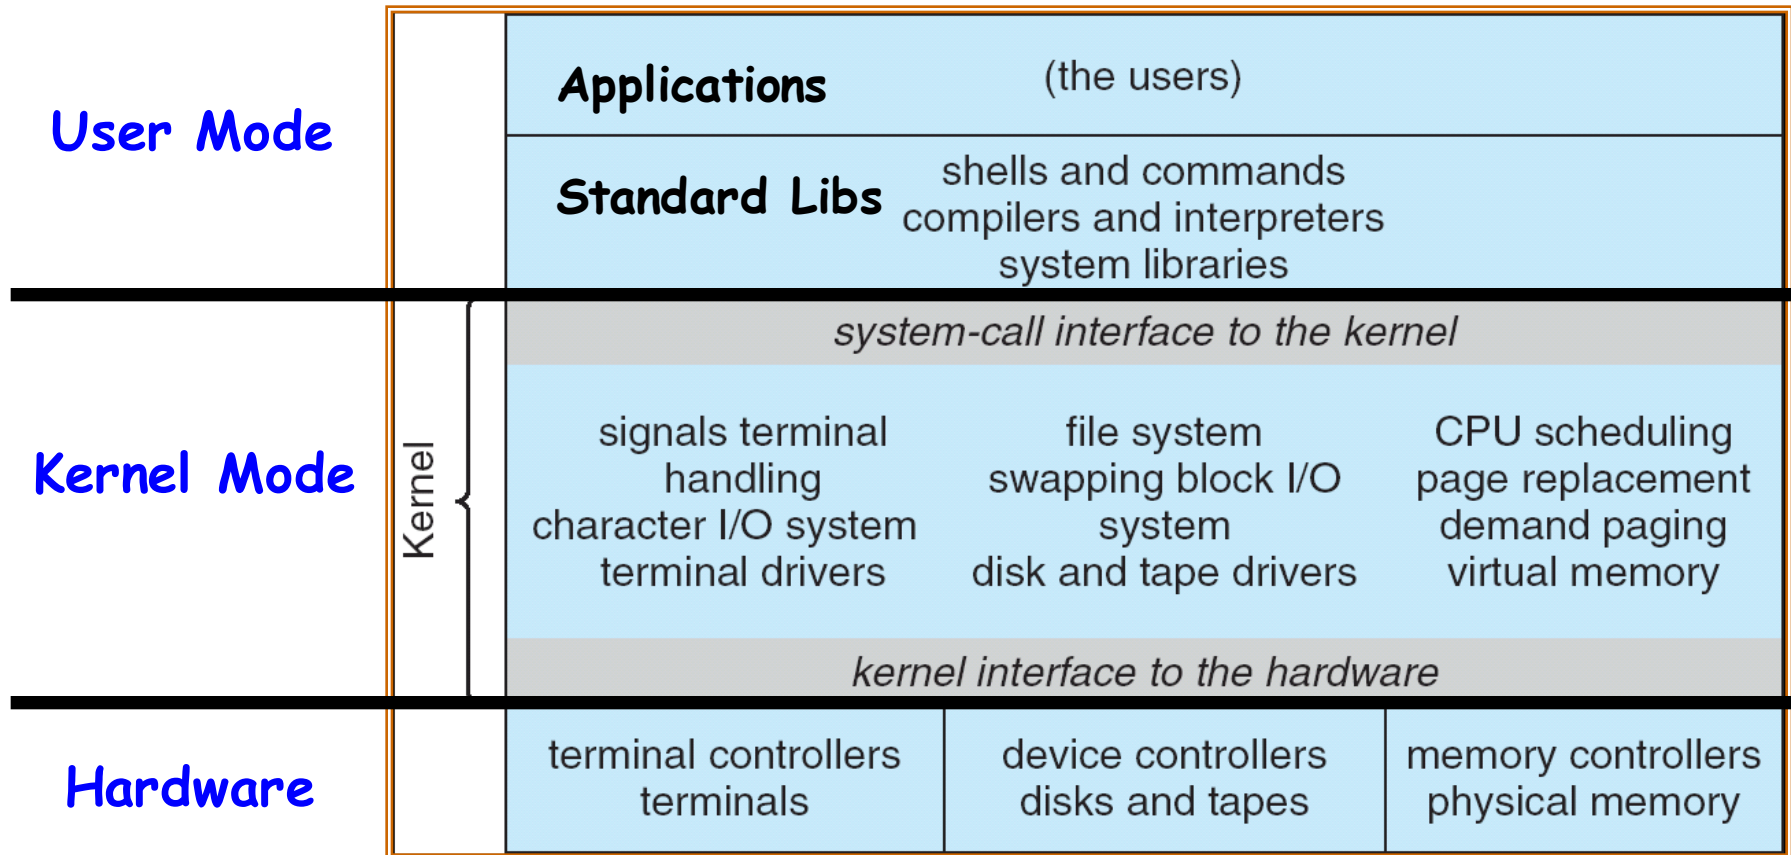
Dual Mode Operation

- **Hardware** provides at least two modes:
 - ❑ “Kernel” mode (or “supervisor” or “protected”)
 - ❑ “User” mode: Normal programs executed
- Some instructions/ops prohibited in user mode:
 - ❑ Example: cannot modify page tables in user mode
 - Attempt to modify \Rightarrow Exception generated
- Transitions from user mode to kernel mode:
 - ❑ System Calls, Interrupts, Other exceptions




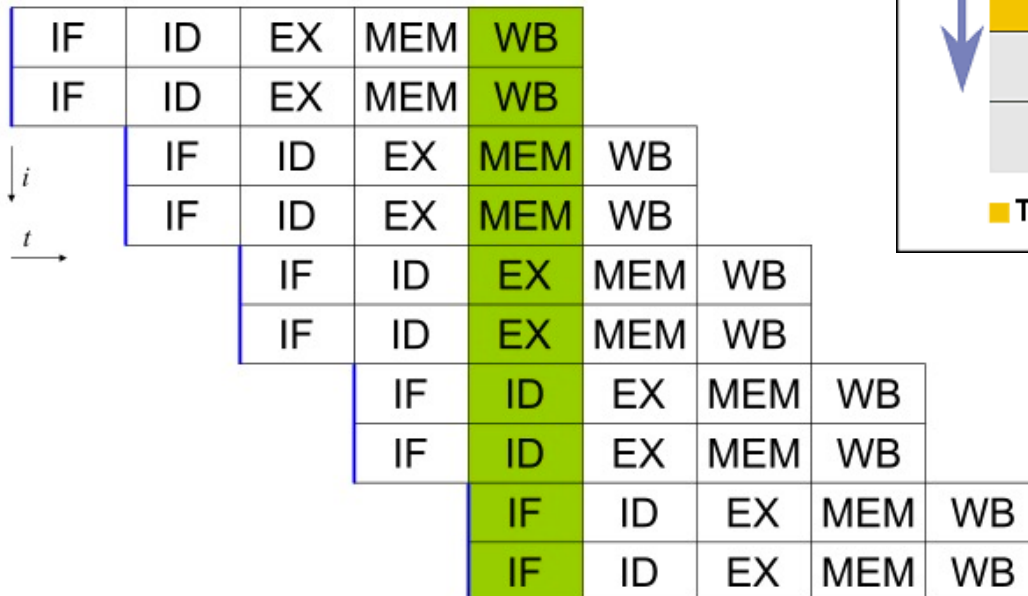
UNIX System Structure

➤ Dual mode: User & Kernel Modes



Modern Technique: SMT/Hyperthreading

- Superscalar processor
 - ❑ instruction-level parallelism within a single processor



Modern Technique: SMT/Hyperthreading

➤ Hardware technique

- ❑ Exploit natural properties of superscalar processors to provide illusion of multiple processors
- ❑ Higher utilization of processor resources

➤ Can schedule each thread as if were separate CPU

- ❑ However, not linear speedup!
- ❑ If have multiprocessor, should schedule each processor first

➤ Original technique called “Simultaneous Multithreading”

- ❑ See <http://www.cs.washington.edu/research/smt/>
- ❑ Alpha, SPARC, Pentium 4 (“Hyperthreading”), Power 5

