

CS 630 - 002

Operating Systems Design

Lecture 3

CPU Scheduling

Jing Li

jingli@njit.edu

GITC 4106



New Jersey Institute of Technology

In This Lecture

➤ CPU Scheduling

- ❑ Scheduling policy goals
- ❑ Scheduling policy options
- ❑ Practical considerations for implementations

Slides courtesy of Hung Daochuan, Chris Gill, David Ferry, Tarek Abdelzaher, Ion Stoica, John Kubiawicz, Peter Dennings, Anthony Joseph, Jonathan Ragan-Kelley, Peter Troger, Insup Lee.

Recap




New Jersey Institute of Technology


Different Operating Systems Designs

- Simple:
 - ❑ Only one or two levels of code
- Layered:
 - ❑ Lower levels independent of upper levels
- Microkernel:
 - ❑ OS built from many user-level processes
- Modular:
 - ❑ Core kernel with Dynamically loadable modules
- ExoKernel:
 - ❑ Separate protection from management of resources
- Cell-based OS:
 - ❑ Two-level scheduling of resources

Recall: What happens during program execution?



How to give the illusion of multiple processors?



- Assume a single processor. How do we provide the illusion of multiple processors?
 - ❑ Multiplex in time!
- Each virtual “CPU” needs a structure to hold:
 - ❑ Program Counter (PC), Stack Pointer (SP)
 - ❑ Registers (Integer, Floating point, others...?)
 - ❑ **Call result a “Thread” for now...**
- How to switch from one CPU to the next?
 - ❑ Save PC, SP, and registers in current state block
 - ❑ Load PC, SP, and registers from new state block
- What triggers switch?
 - ❑ Timer, voluntary yield, I/O, other things

Protecting Threads from Each Other

- Problem: Run multiple applications in such a way that they are protected from one another
- Goal:
 - ❑ Keep User Programs from Crashing OS
 - ❑ Keep User Programs from Crashing each other
 - ❑ [Keep Parts of OS from crashing other parts?]
- (Some of the required) Mechanisms:
 - ❑ Address Translation
 - ❑ Dual Mode Operation
- Simple Policy:
 - ❑ Programs are not allowed to read/write memory of other Programs or of Operating System

Modern Technique: SMT/Hyperthreading

➤ Hardware technique

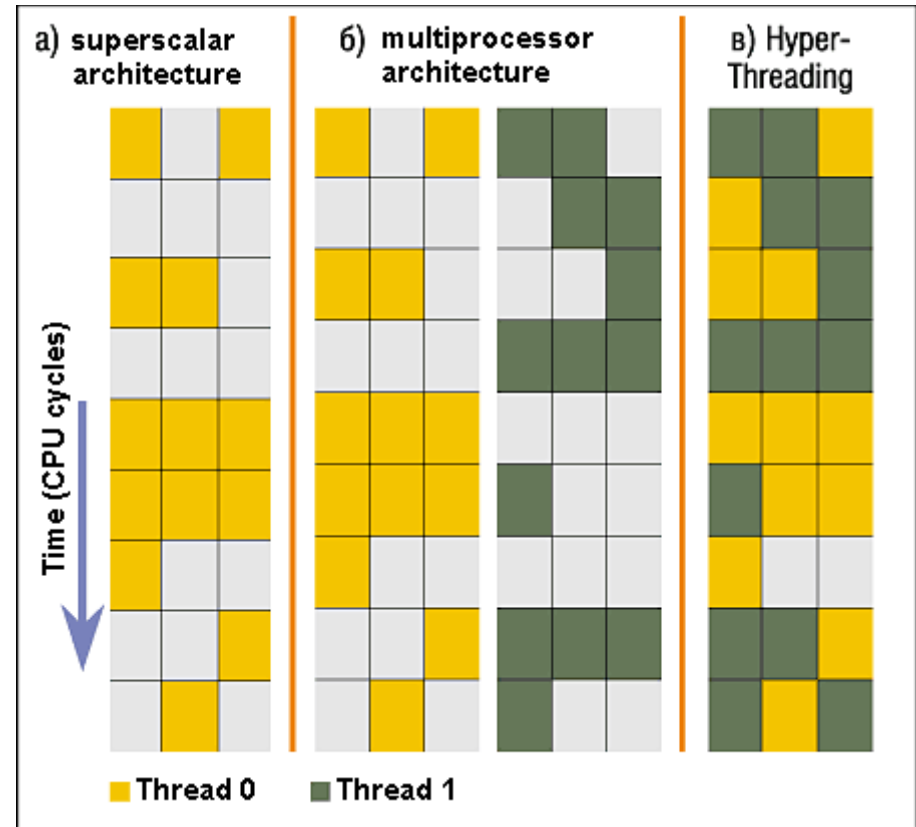
- ❑ Exploit natural properties of superscalar processors to provide illusion of multiple processors
- ❑ Higher utilization of processor resources

➤ Can schedule each thread as if were separate CPU

- ❑ However, not linear speedup!
- ❑ If have multiprocessor, should schedule each processor first

➤ Original technique called “Simultaneous Multithreading”

- ❑ See <http://www.cs.washington.edu/research/smt/>
- ❑ Alpha, SPARC, Pentium 4 (“Hyperthreading”), Power 5

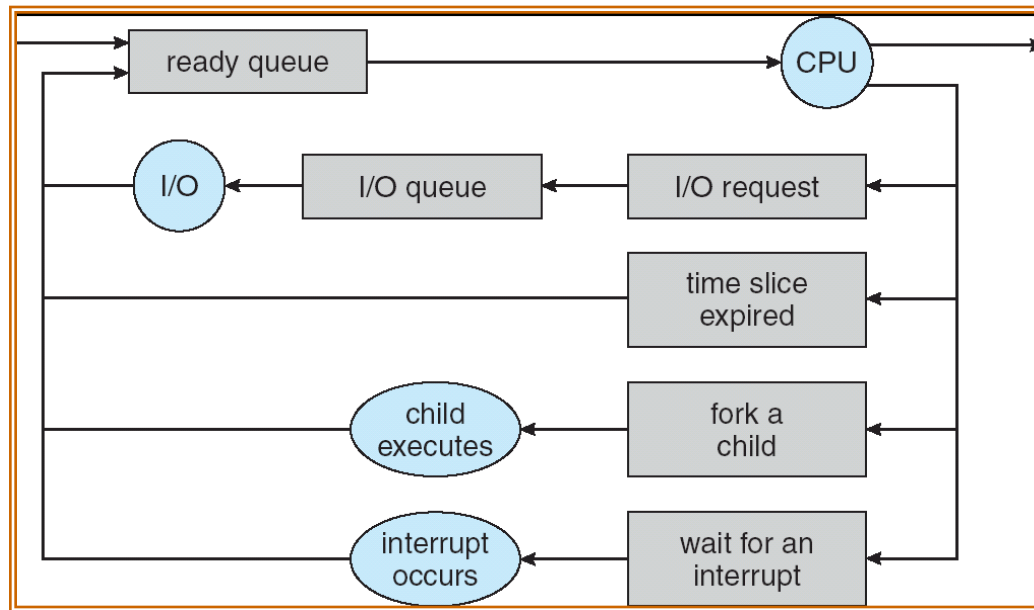


Scheduling



New Jersey Institute of Technology

CPU Scheduling



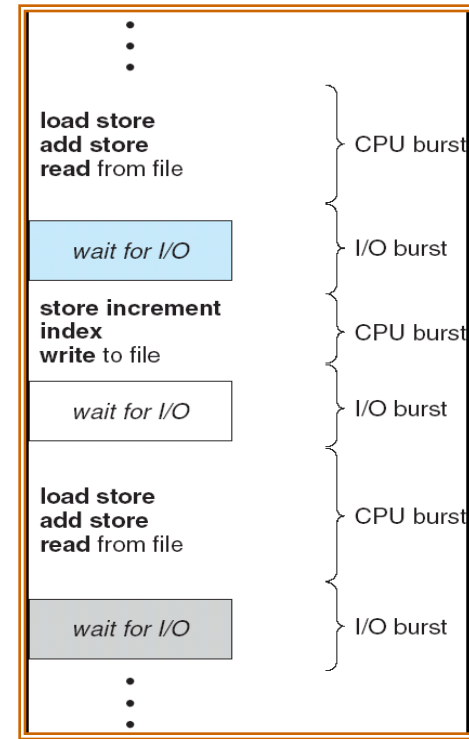

- Question: How is the OS to decide which of several tasks to take off a queue?
 - ❑ Obvious queue to worry about is ready queue
 - ❑ Others can be scheduled as well
- **Scheduling**: deciding which threads are given access to resources from moment to moment

(Simplified) Scheduling Assumptions

- CPU scheduling big area of research in early 70's
- Many implicit assumptions for CPU scheduling:
 - ❑ One program per user
 - ❑ One thread per program
 - ❑ Programs are independent
- Clearly, these are unrealistic but they simplify the problem so it can be solved
 - ❑ For instance: is “fair” about fairness among users or programs?
 - If I run one compilation job and you run five, you get five times as much CPU on many operating systems
- The high-level goal: Dole out CPU time to optimize some desired parameters of system



Assumption: CPU Bursts



- Execution model: programs alternate between CPU & I/O bursts
 - ❑ Program typically uses the CPU for some period of time, then does I/O, then uses CPU again
 - ❑ Each scheduling decision is about which job to give to the CPU for use by its next CPU burst
 - ❑ With time slicing, thread may be forced to give up CPU before finishing current CPU burst
- CPU burst distribution is of main concern

Scheduling Policy Goals/Criteria

➤ Minimize Response Time

- ❑ Minimize elapsed time to do an operation (or job)
- ❑ Response time is what the user sees:
 - Time to echo a keystroke in editor
 - Time to compile a program
 - Real-time Tasks: Must meet deadlines imposed by World

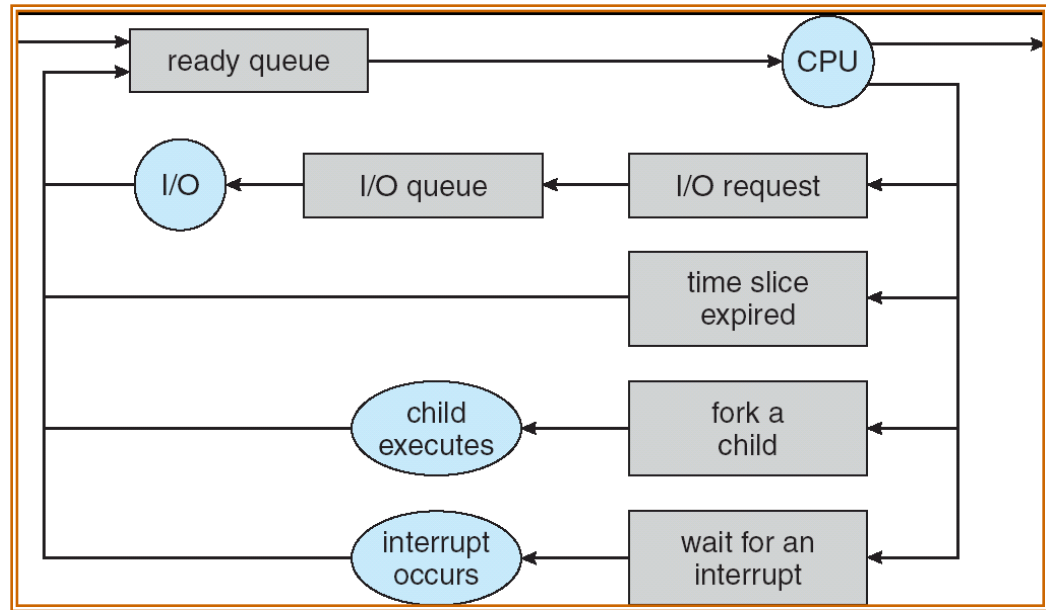
➤ Maximize Throughput

- ❑ Maximize operations (or jobs) per second
- ❑ Throughput related to response time, but not identical:
 - Minimizing response time will lead to more context switching than if you only maximized throughput
- ❑ Two parts to maximizing throughput
 - Minimize overhead (for example, context-switching)
 - Efficient use of resources (CPU, disk, memory, etc)

➤ Fairness

- ❑ Share CPU among users in some equitable way
- ❑ Fairness is not minimizing average response time:
 - Better *average* response time by making system *less* fair

CPU Scheduler



- CPU scheduling decisions may take place when a process:
 - ❑ 1. Switches from running to waiting state
 - ❑ 2. Switches from running to ready state
 - ❑ 3. Switches from waiting to ready (i.e. by an interrupt)
 - ❑ 4. Terminates
- Scheduling under 1 and 4 is **nonpreemptive**
- All other scheduling is **preemptive**

First-Come First-Served (FCFS) Scheduling

➤ First-Come First-Served (FCFS)


❑ Also “First In First Out” (FIFO) or “Run until done”

- In early systems, FCFS meant one program scheduled until done (including I/O)
- Now, means keep CPU until thread blocks

➤ Example:

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

❑ Suppose processes arrive in the order: P_1, P_2, P_3
The schedule is:




❑ Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$

❑ Average waiting time: $(0 + 24 + 27)/3 = 17$

❑ Average response time: $(24 + 27 + 30)/3 = 27$


➤ Convoy effect: short process behind long process



FCFS Scheduling (Cont.)

➤ Example continued:

- ❑ Suppose that processes arrive in order: P_2, P_3, P_1



- ❑ Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
 - ❑ Average waiting time: $(6 + 0 + 3)/3 = 3$
 - ❑ Average response time: $(3 + 6 + 30)/3 = 13$
- ### ➤ In second case:
- ❑ Average waiting time is much better (before it was 17)
 - ❑ Average response time is better (before it was 27)
- ### ➤ FIFO Pros and Cons:
- ❑ Simple (+)
 - ❑ Short jobs get stuck behind long ones (-)

Round Robin (RR)

➤ FCFS Scheme: Potentially bad for short jobs!

- ❑ Depends on submission order
- ❑ If you are first in line at supermarket with milk, you don't care who is behind you, on the other hand...

➤ Round Robin (RR) Scheme

- ❑ Each process gets a small unit of CPU time (*time quantum*), usually 10-100 milliseconds
- ❑ After quantum expires, the process is *preempted* and added to the end of the ready queue.
- ❑ n processes in ready queue and time quantum is $q \Rightarrow$
 - Each process gets $1/n$ of the CPU time
 - In chunks of at most q time units
 - No process waits more than $(n-1)q$ time units




Example of RR with Time Quantum = 20

➤ Example:

<u>Process</u>	<u>Burst Time</u>
P_1	53
P_2	8
P_3	68
P_4	24

❑ The schedule is:



❑ Waiting time for $P_1 = (68-20) + (112-88) = 72$ $P_2 = (20-0) = 20$

$$P_3 = (28-0) + (88-48) + (125-108) = 85$$

$$P_4 = (48-0) + (108-68) = 88$$


❑ Average waiting time = $(72+20+85+88)/4 = 66\frac{1}{4}$

❑ Average response time = $(125+28+153+112)/4 = 104\frac{1}{2}$

Round-Robin Discussion

➤ How do you choose time slice?

- ❑ What if too big?
 - Response time suffers
- ❑ What if infinite (∞)?
 - Get back to FCFS
- ❑ What if time slice too small?
 - Throughput suffers!



➤ Actual choices of timeslice:

- ❑ Initially, UNIX timeslice one second:
 - Worked ok when UNIX was used by one or two people.
 - What if three compilations going on? 3 seconds to echo each keystroke!
- ❑ In practice, need to balance short-job performance and long-job throughput:
 - Typical time slice today is between 10ms – 100ms
 - Typical context-switching overhead is 0.1ms – 1ms
 - Roughly 1% overhead due to context-switching

Comparisons between FCFS and Round Robin

- Assuming zero-cost context-switching time, is RR always better than FCFS?
- Simple example:
- Response Times:

10 jobs, each take 100s of CPU time

RR scheduler quantum of 1s

All jobs start at the same time

Job #	FCFS	RR
1	100	991
2	200	992
...
9	900	999
10	1000	1000

- ❑ Both RR and FCFS finish at the same time
- ❑ Average response time is much worse under RR!
 - RR is bad when all jobs have the same length
- Also: Cache state must be shared between all jobs with RR but can be devoted to each job with FCFS
- ❑ Total time for RR longer even for zero-cost switch!

Break

45 min. Break 45 min. Break 50 min. End

Time 



New Jersey Institute of Technology

Administrivia



New Jersey Institute of Technology

Lab 1: Hello, Linux!



New Jersey Institute of Technology

Purpose of Lab 1

- Will be released tonight and will be **due on Feb 19 23:55**
- Your first trial of forming your group
- Ease you into programming in Linux
- Bases for a later Lab – writing a shell program

In this lab assignment, you will

- 1. Learn some basic UNIX/Linux commands
- 2. Use the built-in Linux manual (man pages) to look up certain functions and commands
- 3. Try to learn and use an editor to create and write files
- 4. Compile a "Hello, Linux!" program using the C standard library

Requirements

- 1. Linux machine (or a MacOS, but not Windows)
- 2. following the exact order
- 3. record your answers by taking a screenshot

```
yuchen@yuchen-VirtualBox:/bin$ ls |head -1  
bash  
yuchen@yuchen-VirtualBox:/bin$
```

- 4. submit your report in the PDF form to Canvas
 - ❑ Late penalty: see first lecture slides or syllabus
 - ❑ Only person in the team need to submit

Earlier Example with Different Time Quantum

Best FCFS:

P ₂ [8]	P ₄ [24]	P ₁ [53]	P ₃ [68]
-----------------------	------------------------	------------------------	------------------------

0 8 32 85 153

	Quantum	P ₁	P ₂	P ₃	P ₄	Average
Wait Time	Best FCFS	32	0	85	8	31¼
	Q = 1	84	22	85	57	62
	Q = 5	82	20	85	58	61¼
	Q = 8	80	8	85	56	57¼
	Q = 10	82	10	85	68	61¼
	Q = 20	72	20	85	88	66¼
	Worst FCFS	68	145	0	121	83½
Response Time	Best FCFS	85	8	153	32	69½
	Q = 1	137	30	153	81	100½
	Q = 5	135	28	153	82	99½
	Q = 8	133	16	153	80	95½
	Q = 10	135	18	153	92	99½
	Q = 20	125	28	153	112	104½
	Worst FCFS	121	153	68	145	121¾

What if we Knew the Future?

- Could we always mirror best FCFS?
 - ❑ Idea is to get short jobs out of the system
 - ❑ Big effect on short jobs, only small effect on long ones
 - ❑ Result is better average response time
- (Non-preemptive) Shortest Job First (SJF, Shortest Process Next):
 - ❑ Run whatever job has the least amount of computation to do
- Shortest Remaining Time First (SRTF, SRPT):
 - ❑ If job arrives and has a shorter time to completion than the remaining time on the current job, immediately preempt CPU
- What is the difference between (non-preempt) SJF and SRTF?
 - ❑ What if when jobs arrive at the same time?
 - ❑ What if when a short job arrives before a long job's completion?




Difference between SJF and SRTF

➤ Example:


Process	Arrival Time	Burst Time
P_1	0	20
P_2	10	7
P_3	26	3

❑ (Non-Preemptive) SJF:




Average response time:
 $(20 + 17 + 4)/3 = 13.67$

❑ SRTF



Average response time:
 $(27 + 7 + 4)/3 = 12.67$

❑ (Preemptive SJF: strictly speaking, different from SRTF)




Average response time:
 $(30 + 7 + 3)/3 = 13.3$

Discussion

- SJF/SRTF are the best for minimizing average response time
 - ❑ Provably optimal
 - (SJF among non-preemptive, SRTF among preemptive)
 - ❑ Since SRTF is always at least as good as SJF, focus on SRTF
- Comparison of SRTF with FCFS and RR
 - ❑ What if all jobs the same length?
 - SRTF becomes the same as FCFS
 - (i.e. FCFS is the best one can do if all jobs have the same length)
 - ❑ What if jobs have varying length?
 - SRTF (and RR): short jobs not stuck behind long ones

Example to illustrate benefits of SRTF



➤ Three jobs:

- ❑ A,B: both CPU bound, run for a week
C: I/O bound, loop 1ms CPU, 9ms disk I/O
- ❑ If only one at a time, C uses 90% of the disk, A or B could use 100% of the CPU

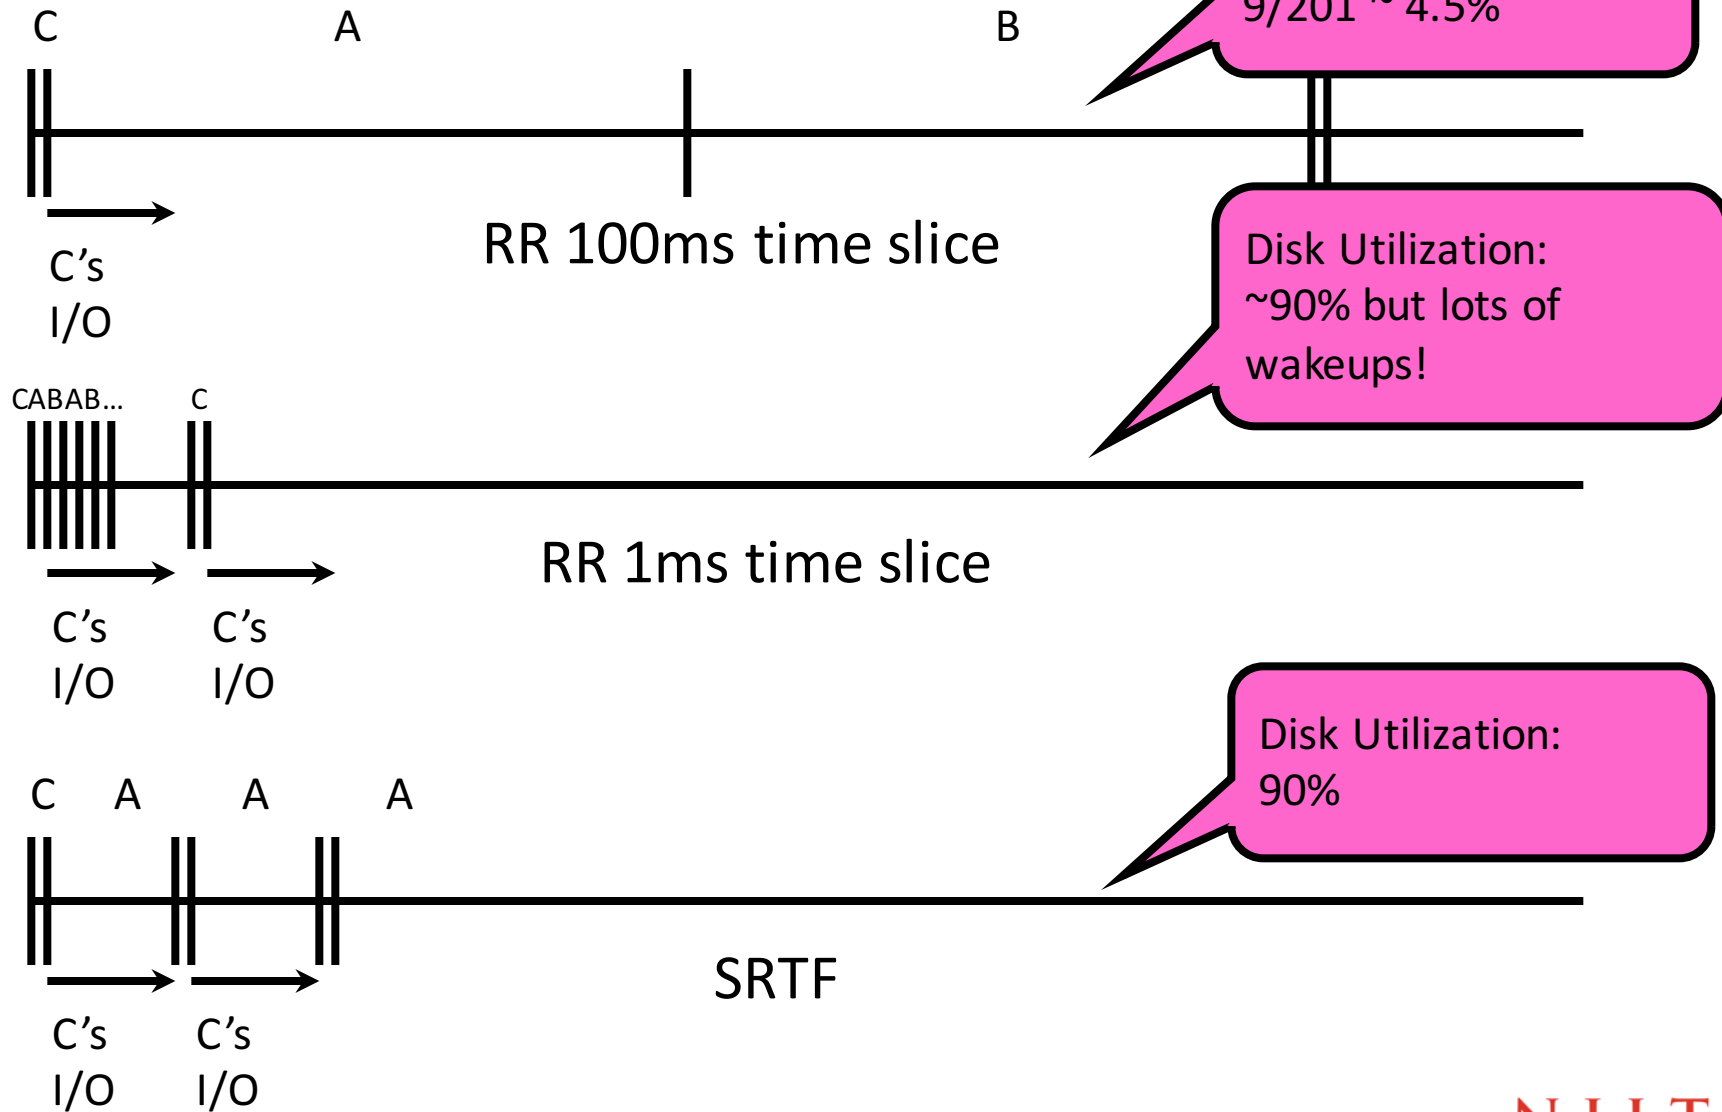
➤ With FCFS:

- ❑ Once A or B get in, keep CPU for two weeks

➤ What about RR or SRTF?

- ❑ Easier to see with a timeline

SRTF Example continued:



SRTF Further discussion

➤ Starvation

- ❑ Large jobs never get to run (starvation), if many small jobs

➤ Somehow need to predict future

- ❑ How can we do this?
- ❑ Some systems ask the user
 - When you submit a job, have to say how long it will take
 - To stop cheating, system kills job if takes too long
- ❑ But: Even non-malicious users have trouble predicting runtime of their jobs

➤ Bottom line, can't really know how long job will take

- ❑ However, can use SRTF as a yardstick for measuring other policies
- ❑ Optimal, so can't do any better

➤ SRTF Pros & Cons

- ❑ Optimal (average response time) (+)
- ❑ Hard to predict future (-)
- ❑ Unfair (-)



Summary

- **Scheduling:** selecting a waiting process from the ready queue and allocating the CPU to it
- **FCFS Scheduling:**
 - ❑ Run jobs to completion in order of submission
 - ❑ Pros: Simple (+)
 - ❑ Cons: Short jobs get stuck behind long ones (-)
- **Round-Robin Scheduling:**
 - ❑ Give each thread a small amount of CPU time when it executes; cycle between all ready threads
 - ❑ Pros: Better for short jobs (+)
 - ❑ Cons: Poor when jobs are same length (-)
- **Shortest Remaining Time First (SRTF):**
 - ❑ Run whatever job has the least amount of computation to do/least remaining amount of computation to do
 - ❑ Pros: Optimal (average response time) (+)
 - ❑ Cons: Hard to predict future, Unfair (-)